ETH zürich

ALMA MATER STUDIORUM
UNIVERSITÀ DI BOLOGNA

SoftHier Progress Update

Chi Zhang chizhang@iis.ee.ethz.ch

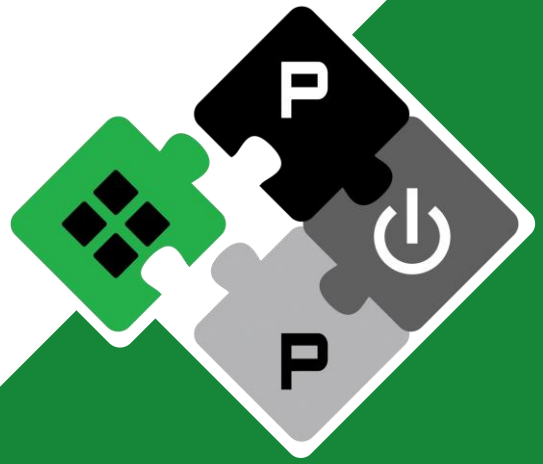

@pulp_platform

pulp-platform.org

youtube.com/pulp_platform

PULP Platform

Open Source Hardware, the way it should be!

SoftHier: Parameterizable NoC-Based Scalable System

- Features
  - Two separate NoC bus
    - DATA NoC: wide link, transfer bulk data
    - SYNC NoC: narrow link, cluster synchronization
  - Infinite Instruction memory
    - Ignore I$ fetch overhead
- Fully Parameterizable
  - Configure file and push button
    - #Cores, RedMule config
    - L1 (TCDM) size & BW
    - #Clusters (#row, #col)
    - #HBM channels and placement
    - NoC link BW
- SW stack ready

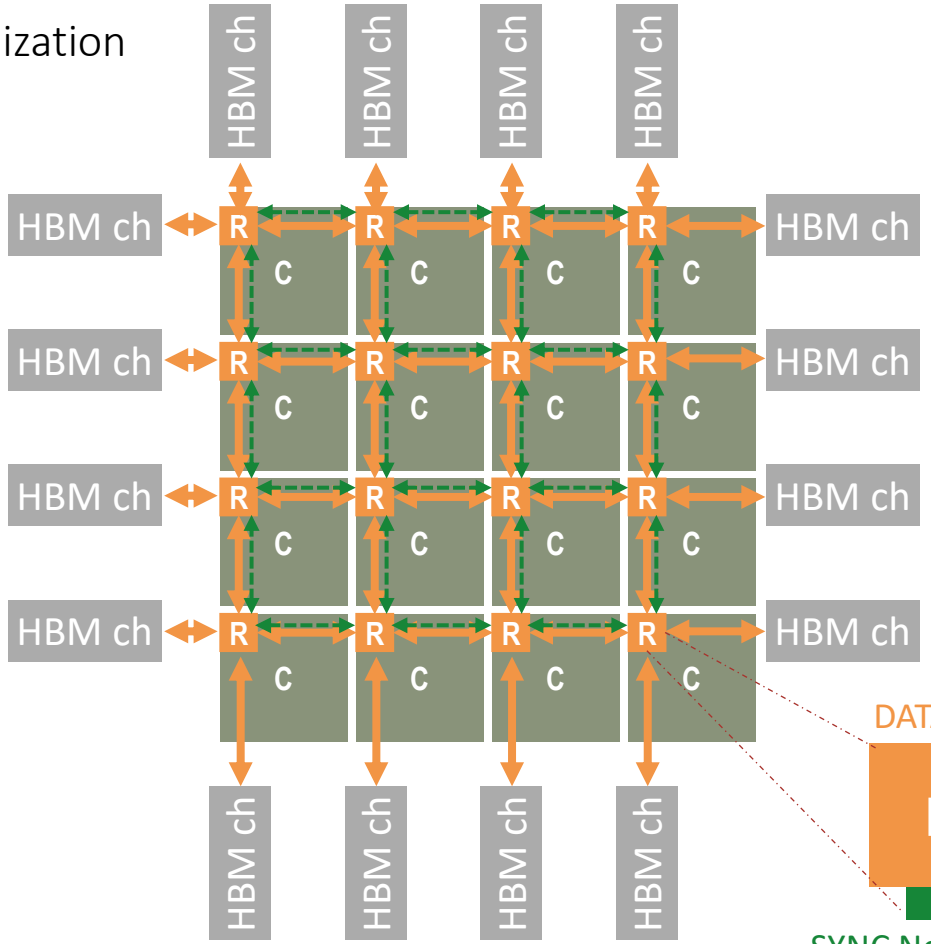

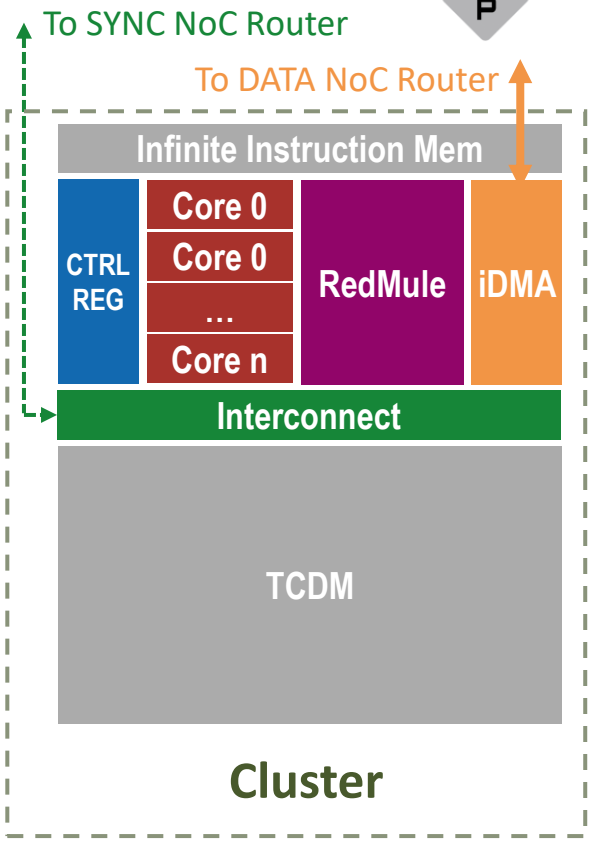

ETH zürich ALMA MATER STUDIORUM UNIVERSITÀ DI BOLOGNA

Design Space Exploration of RedMule

- Goal: We Want to Know
  - What is **optimal MatMul problem size** (M-N-K)to reach the best efficiency of RedMule **in a Cluster:** the **Optimal Efficiency Point**
    - When Matrix too small -> low RedMule utilization
    - When Matrix too large -> we introduce large and redundant TCDM space
    - We are seeking for **best MACs/SRAM in a CLuster**
  - What is the BW needed for RedMule at **Optimal Efficiency Point**
- Design Space Exploration Constraints
  - RedMule CE array constraint: L = H * Pipe, CE Pipe=4
- Key Metric
  - Efficiency Metric:
    - $\frac{Effective\ FLOP/Cycle}{TCDM\ Occupied\ Area} = \frac{RedMule\ Utilization * 2 * \#CEs}{Elem\ Size * 2 * (MN+NK+MK)}$

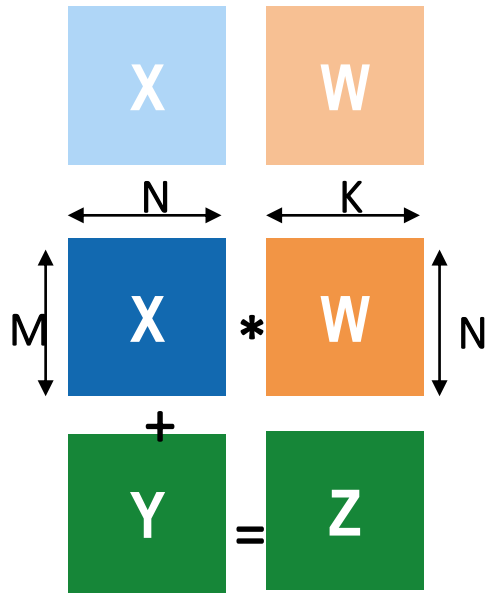

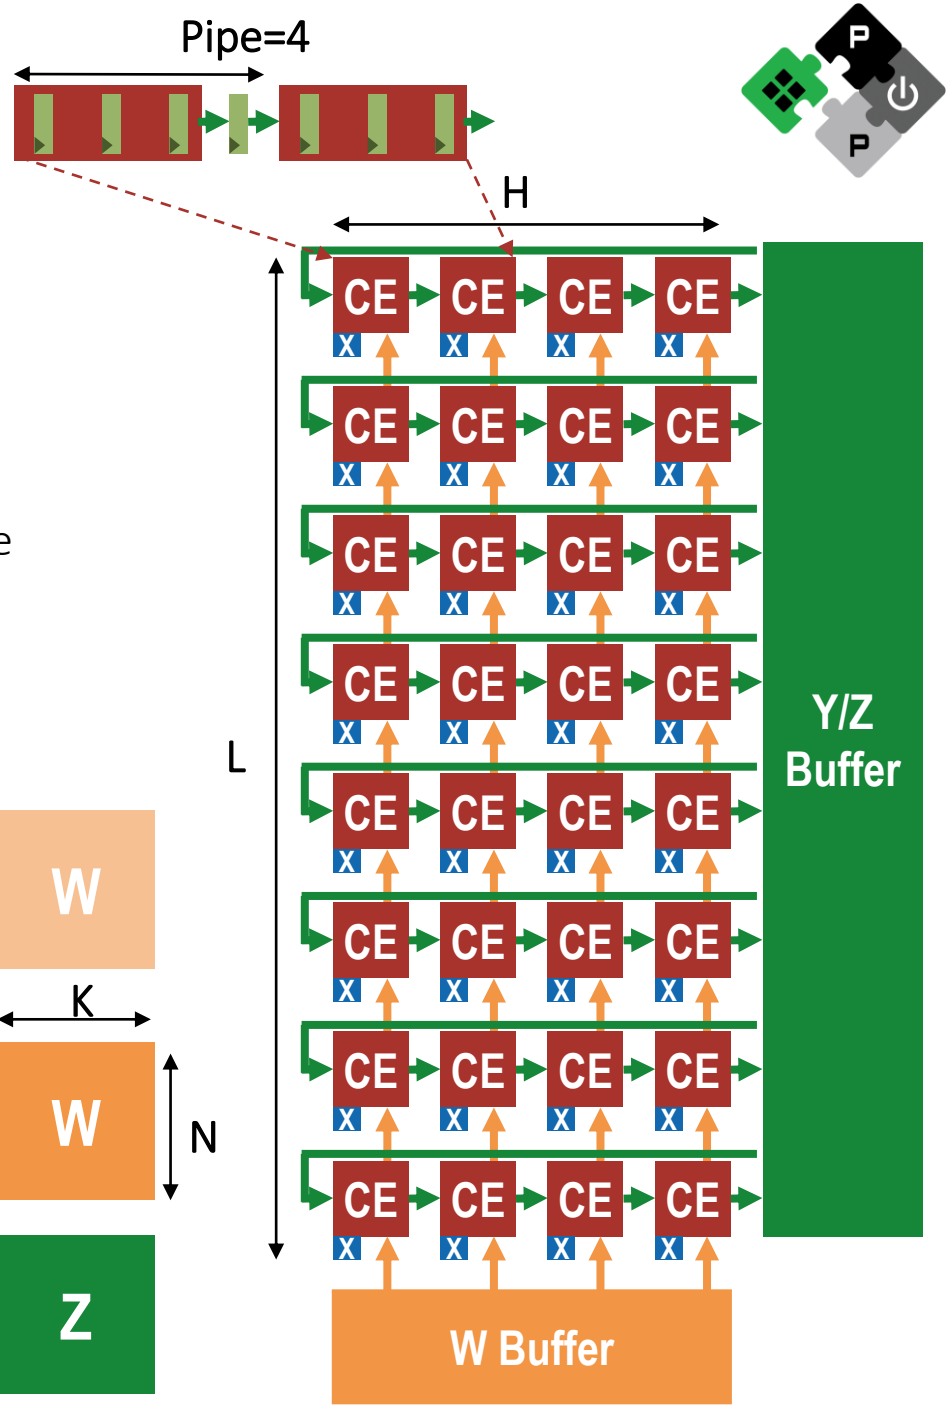

Find Optimal GEMM Size(M-N-K) for RedMule (@ saturated BW)

- GEMM Dimension
  - $M \quad \in [8, 16, 32, 64, 128, 256, 512]$
  - $N \quad \in [8, 16, 32, 64, 128, 256, 512]$
  - $K \quad \in [8, 16, 32, 64, 128, 256, 512]$
- RedMule Config
  - CE array = 128x32
  - TCDM BW = 1024 Elem/Cycle
  - Element = FP16
- Run GEMM on One Cluster with One RedMule
- Collect Data
  - Efficiency Metric:
    - $\frac{Effective\ FLOP/Cycle}{TCDM\ Occupied\ Area} = \frac{RedMule\ Utilization * 2 * \#CEs}{Elem\ Size * 2 * (MN+NK+MK)}$

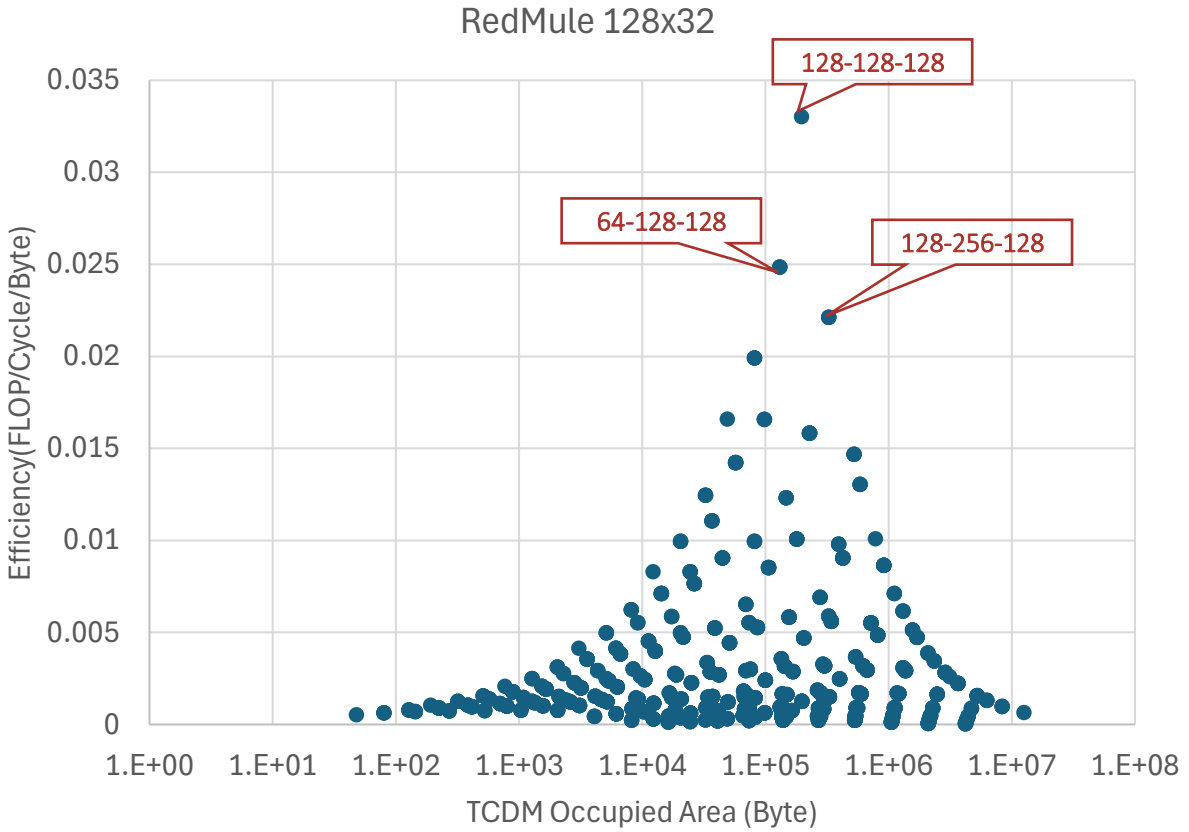

Find Optimal GEMM Size (M-N-K) for RedMule (@ saturated BW)

- At RedMule Constraints: CE array constraint: L = H * Pipe

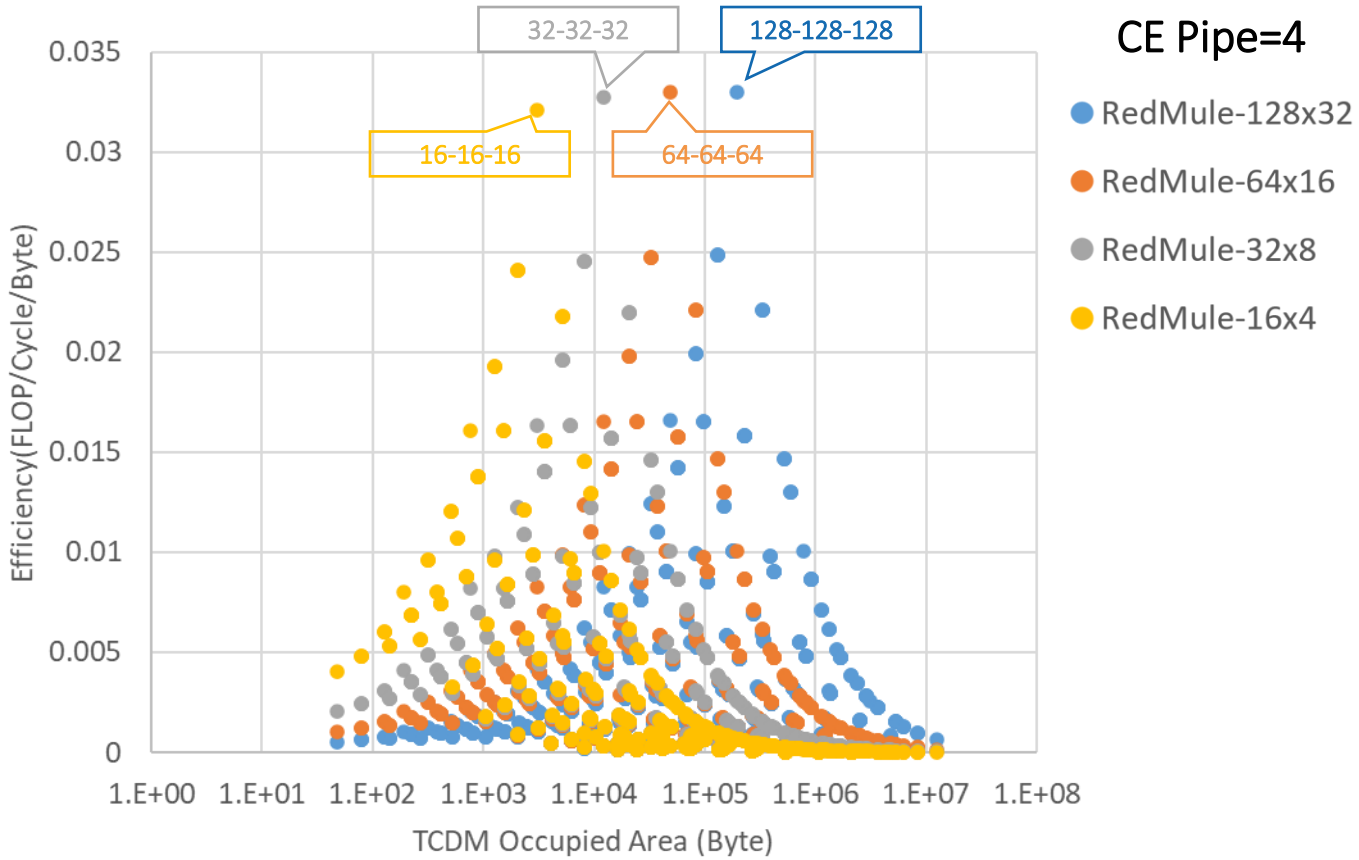

To use RedMule at best efficiency (Optimal Efficiency Point)

$$M = N = K = \sqrt{CE\ array\ *\ CE\ Pipeline} = L$$

Find Optimal GEMM Size (M-N-K) for RedMule (@ saturated BW)

- At RedMule Constraints: CE array constraint: L = H * Pipe

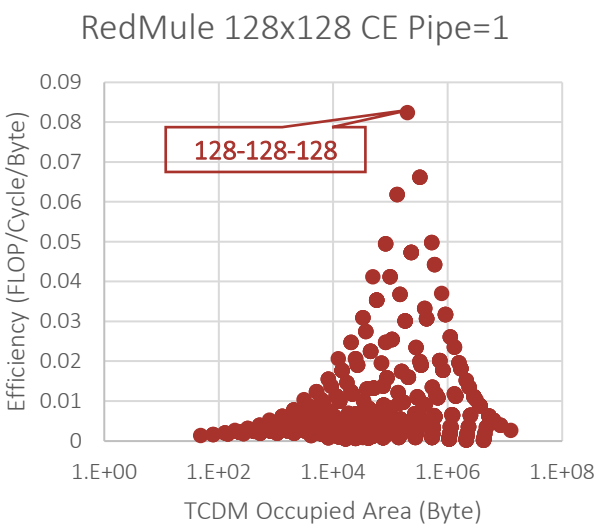

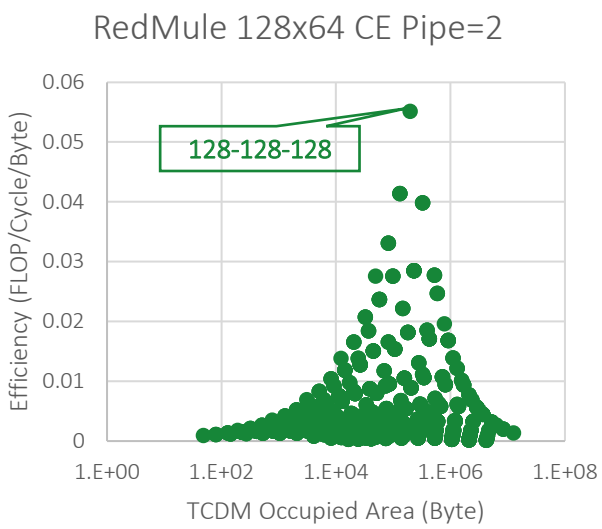

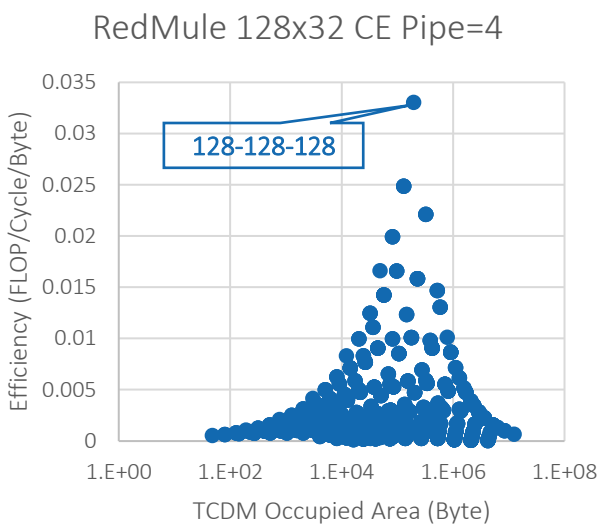

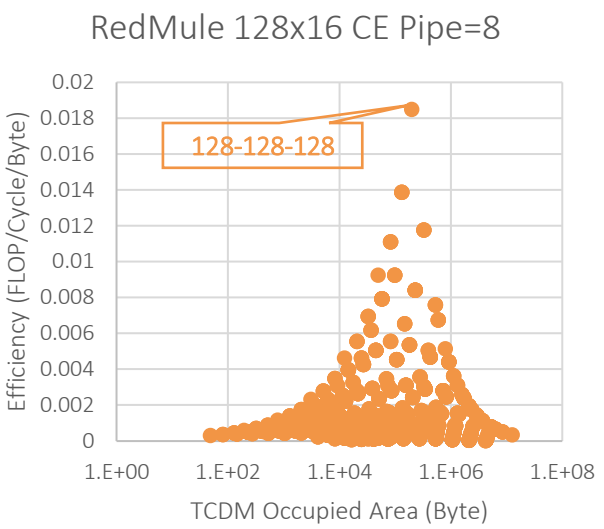

To use RedMule at best efficiency (Optimal Efficiency Point)

$$M = N = K = \sqrt{CE\ array\ *\ CE\ Pipeline} = L$$

TCDM Bandwidth Requirement

- RedMule Utilization Heatmap
  - GEMM Size vs RedMule CE array
  - Vary TCDM BW
  - We're satisfied at > 69% uti

To feed RedMule efficiently TCDM Bandwidth (elem per cycle):

$$BW \geq 4\sqrt{CE\ array\ *\ CE\ Pipeline} = 4L$$

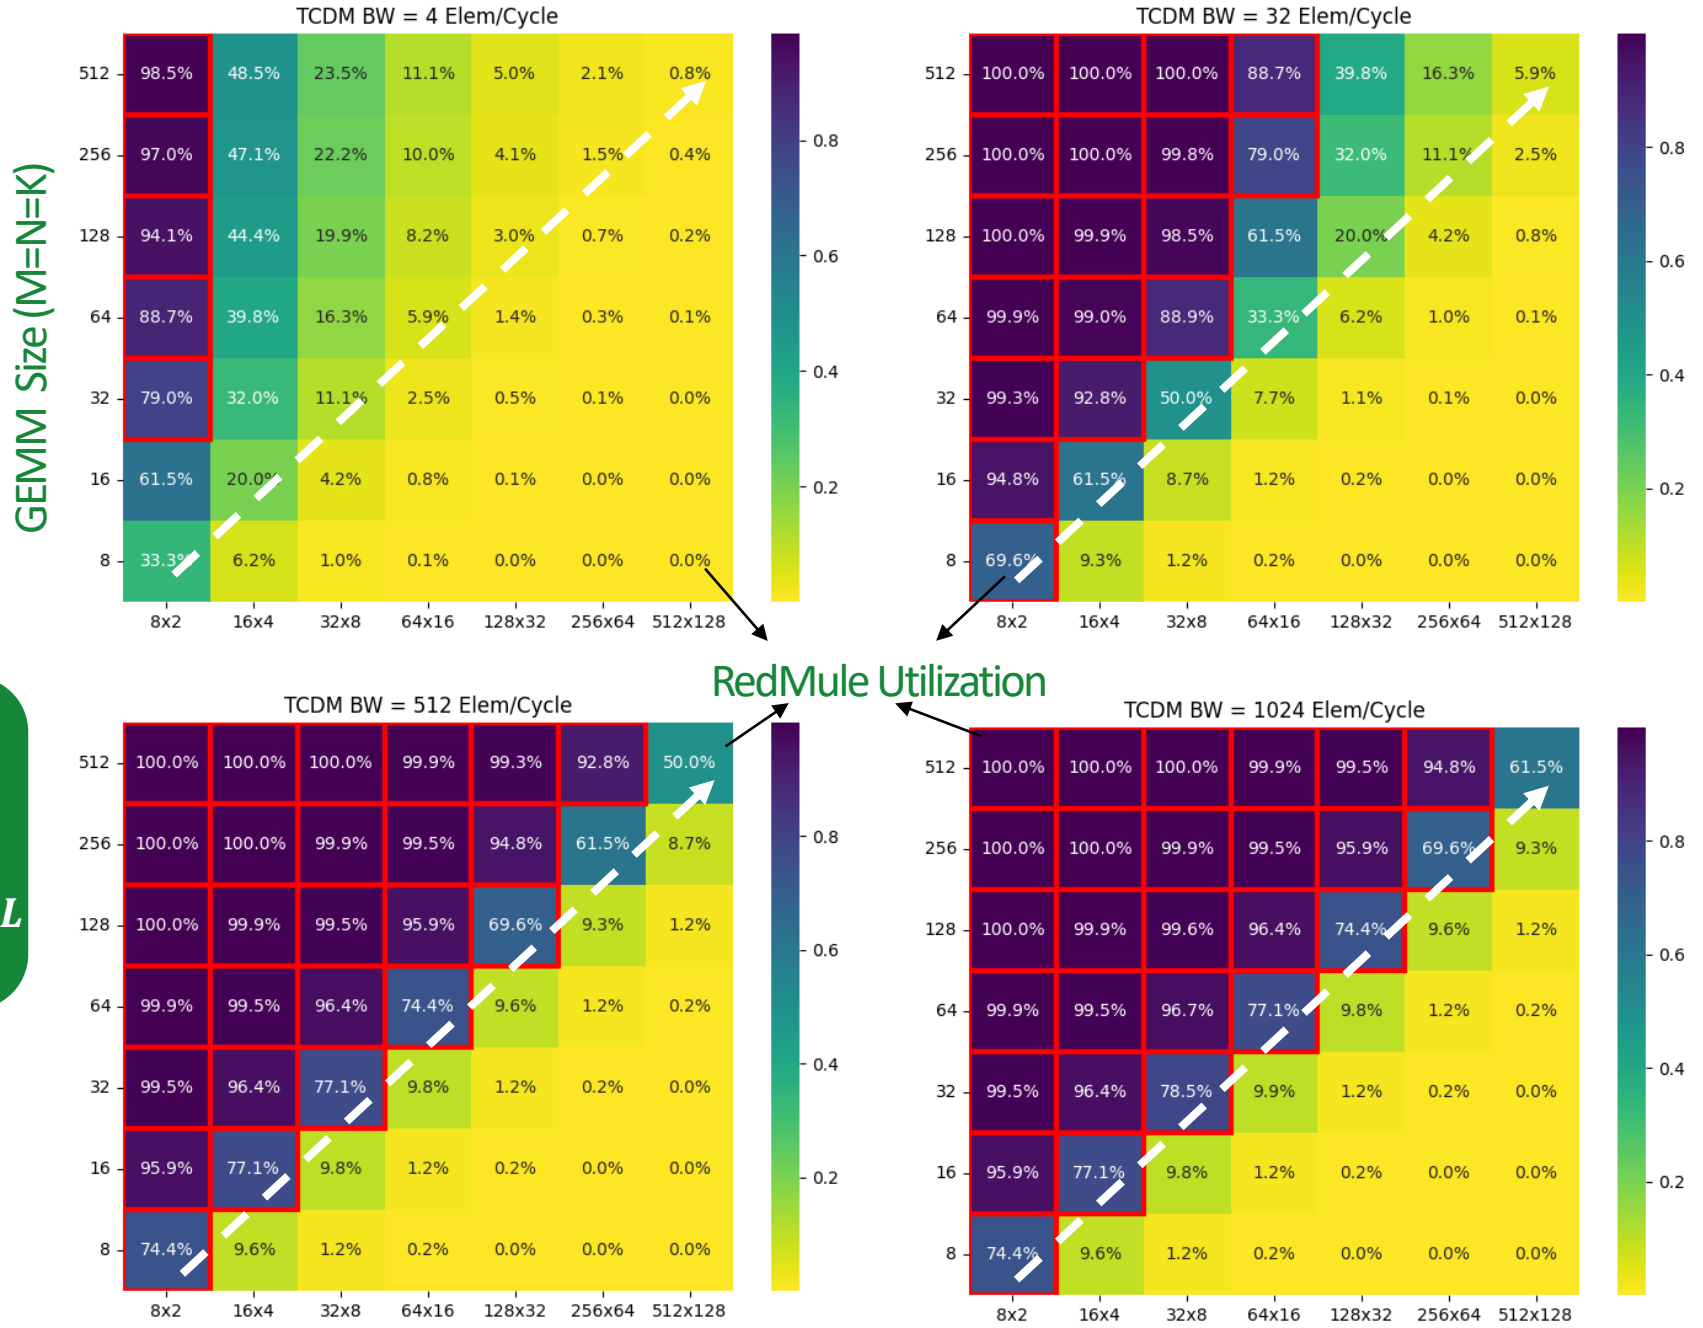

For Large GEMM: Enable Cluster-to-Cluster Comm or Not?

- No Cluster-to-Cluster Comm.
  - Each clusters take care of different output tiles (384x384).
  - iDMA transfers input matrix tile + weight matrix tile from HBM
  - No inter-cluster tile reuses

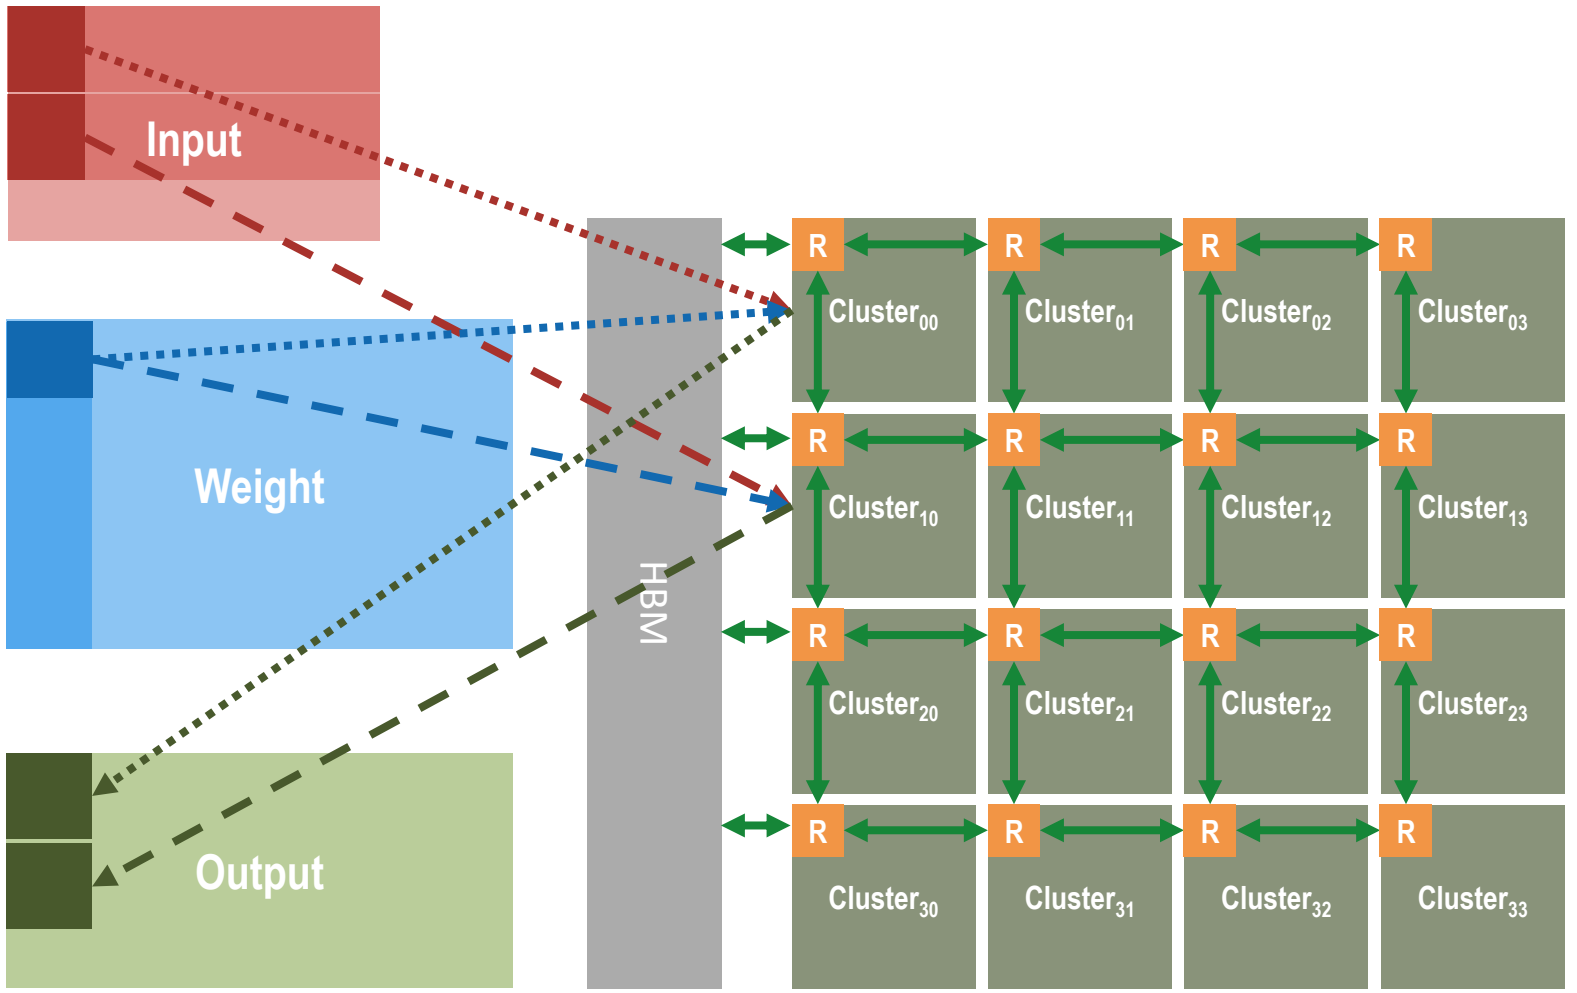

For Large GEMM: Enable Cluster-to-Cluster Comm or Not?

- Leverage Cluster-to-Cluster Comm.
  - Reuse tiles, reduce HBM accesses, save BW limitation
  - Need “smart” tile mapping and scheduling

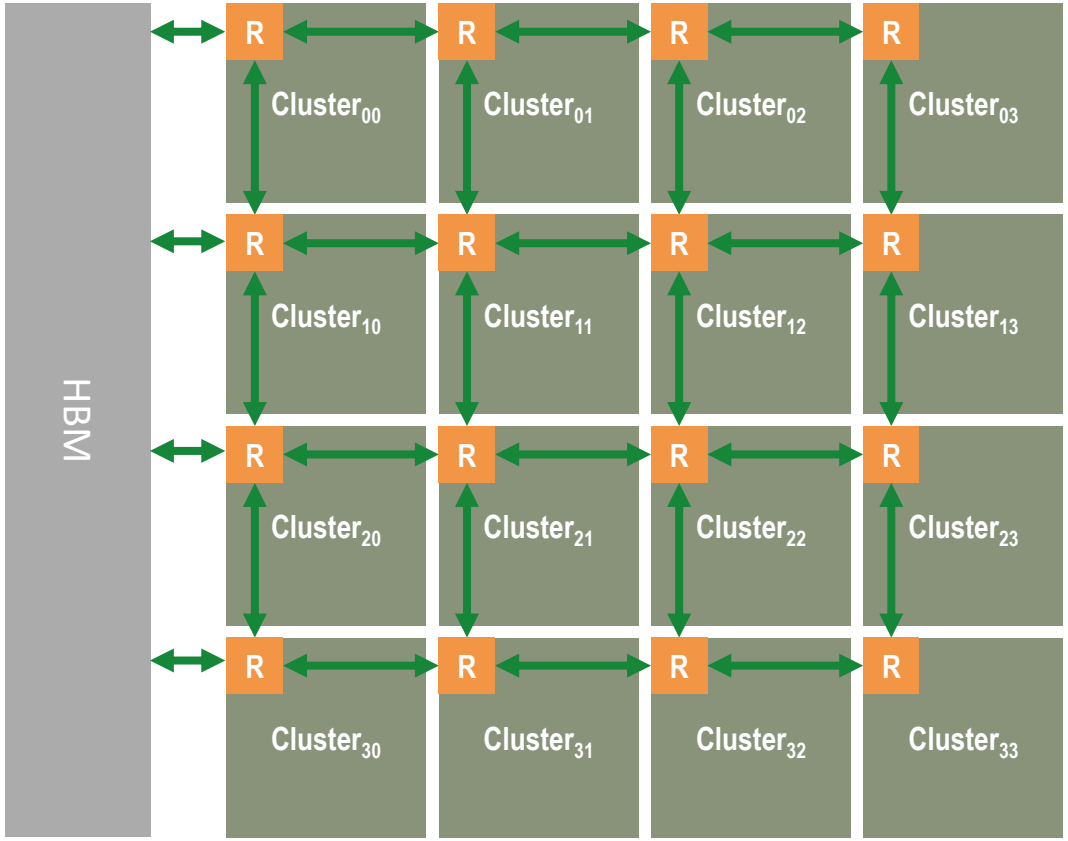

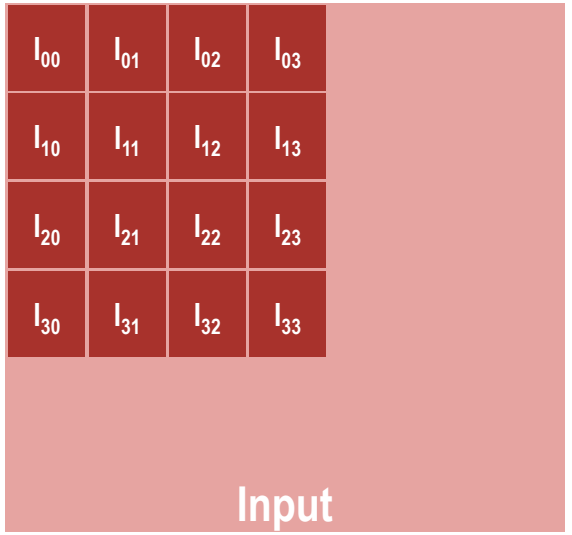

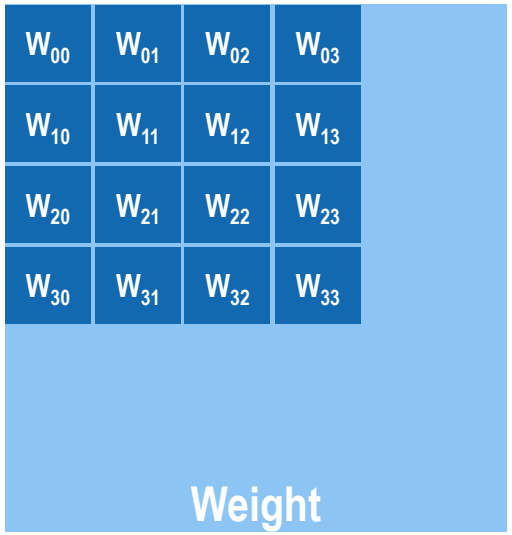

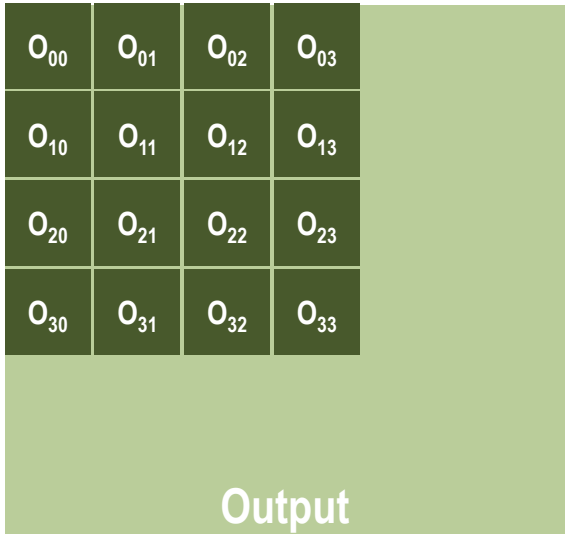

Example Solution for GEMM: Step1 Distribute Tiles to Clusters

- Load All (16 Input tiles, 16 weight tiles) from HBM
- Cluster contain different tiles from each other

$$
\begin{aligned}
O_{00} &= \boxed{I_{00} \cdot W_{00}} + I_{01} \cdot W_{10} + I_{02} \cdot W_{20} + I_{03} \cdot W_{30} \\
O_{01} &= I_{00} \cdot W_{01} + \boxed{I_{01} \cdot W_{11}} + I_{02} \cdot W_{21} + I_{03} \cdot W_{31} \\
O_{02} &= I_{00} \cdot W_{02} + I_{01} \cdot W_{12} + \boxed{I_{02} \cdot W_{22}} + I_{03} \cdot W_{32} \\
O_{03} &= I_{00} \cdot W_{03} + I_{01} \cdot W_{13} + I_{02} \cdot W_{23} + \boxed{I_{03} \cdot W_{33}} \\
O_{10} &= I_{10} \cdot W_{00} + I_{11} \cdot W_{10} + I_{12} \cdot W_{20} + \boxed{I_{13} \cdot W_{30}} \\
O_{11} &= \boxed{I_{10} \cdot W_{01}} + I_{11} \cdot W_{11} + I_{12} \cdot W_{21} + I_{13} \cdot W_{31} \\
O_{12} &= I_{10} \cdot W_{02} + \boxed{I_{11} \cdot W_{12}} + I_{12} \cdot W_{22} + I_{13} \cdot W_{32} \\
O_{13} &= I_{10} \cdot W_{03} + I_{11} \cdot W_{13} + \boxed{I_{12} \cdot W_{23}} + I_{13} \cdot W_{33} \\
O_{20} &= I_{20} \cdot W_{00} + I_{21} \cdot W_{10} + \boxed{I_{22} \cdot W_{20}} + I_{23} \cdot W_{30} \\
O_{21} &= I_{20} \cdot W_{01} + I_{21} \cdot W_{11} + I_{22} \cdot W_{21} + \boxed{I_{23} \cdot W_{31}} \\
O_{22} &= \boxed{I_{20} \cdot W_{02}} + I_{21} \cdot W_{12} + I_{22} \cdot W_{22} + I_{23} \cdot W_{32} \\
O_{23} &= I_{20} \cdot W_{03} + \boxed{I_{21} \cdot W_{13}} + I_{22} \cdot W_{23} + I_{23} \cdot W_{33} \\
O_{30} &= I_{30} \cdot W_{00} + \boxed{I_{31} \cdot W_{10}} + I_{32} \cdot W_{20} + I_{33} \cdot W_{30} \\
O_{31} &= I_{30} \cdot W_{01} + I_{31} \cdot W_{11} + \boxed{I_{32} \cdot W_{21}} + I_{33} \cdot W_{31} \\
O_{32} &= I_{30} \cdot W_{02} + I_{31} \cdot W_{12} + I_{32} \cdot W_{22} + \boxed{I_{33} \cdot W_{32}} \\
O_{33} &= \boxed{I_{30} \cdot W_{03}} + I_{31} \cdot W_{13} + I_{32} \cdot W_{23} + I_{33} \cdot W_{33}
\end{aligned}
$$

Example Solution for GEMM: Step1 Distribute Tiles to Clusters

- Load All (16 Input tiles, 16 weight tiles) from HBM
- Cluster contain different tiles from each other

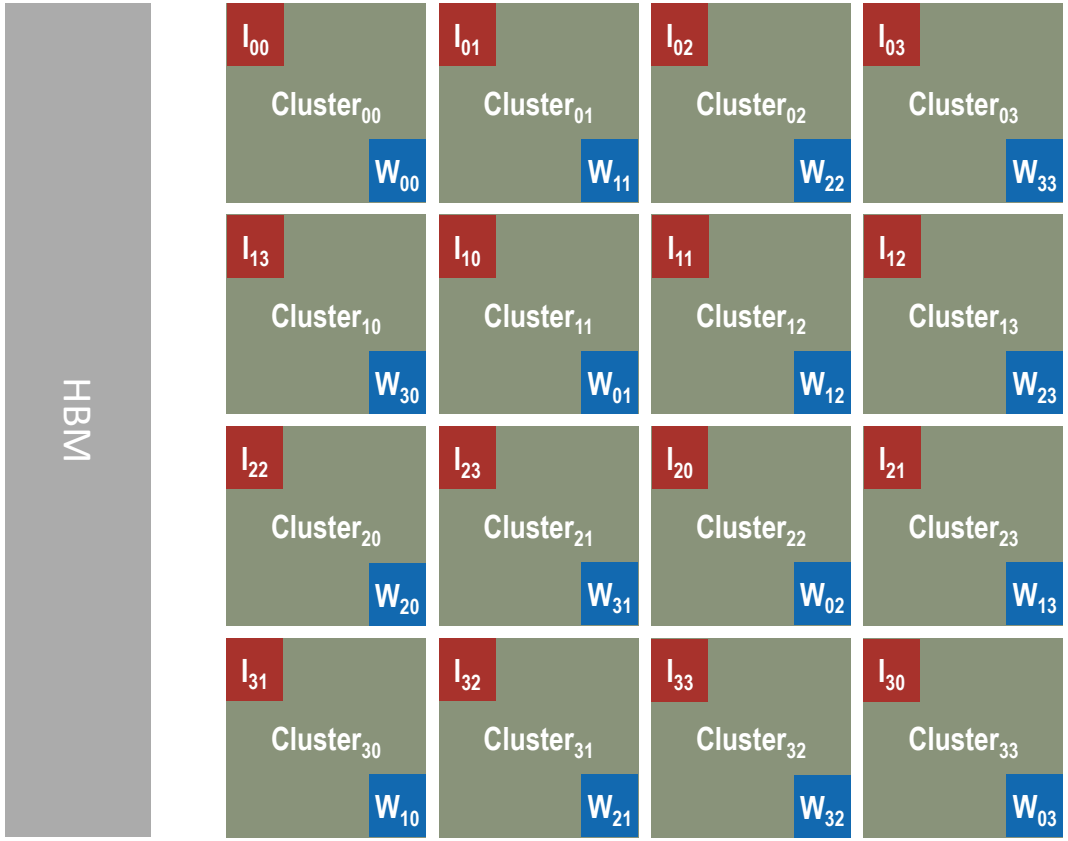

$$
\begin{aligned}
O_{00} &= \boxed{I_{00} \cdot W_{00}} + I_{01} \cdot W_{10} + I_{02} \cdot W_{20} + I_{03} \cdot W_{30} \\
O_{01} &= I_{00} \cdot W_{01} + \boxed{I_{01} \cdot W_{11}} + I_{02} \cdot W_{21} + I_{03} \cdot W_{31} \\
O_{02} &= I_{00} \cdot W_{02} + I_{01} \cdot W_{12} + \boxed{I_{02} \cdot W_{22}} + I_{03} \cdot W_{32} \\
O_{03} &= I_{00} \cdot W_{03} + I_{01} \cdot W_{13} + I_{02} \cdot W_{23} + \boxed{I_{03} \cdot W_{33}} \\
O_{10} &= I_{10} \cdot W_{00} + I_{11} \cdot W_{10} + I_{12} \cdot W_{20} + \boxed{I_{13} \cdot W_{30}} \\
O_{11} &= \boxed{I_{10} \cdot W_{01}} + I_{11} \cdot W_{11} + I_{12} \cdot W_{21} + I_{13} \cdot W_{31} \\
O_{12} &= I_{10} \cdot W_{02} + \boxed{I_{11} \cdot W_{12}} + I_{12} \cdot W_{22} + I_{13} \cdot W_{32} \\
O_{13} &= I_{10} \cdot W_{03} + I_{11} \cdot W_{13} + \boxed{I_{12} \cdot W_{23}} + I_{13} \cdot W_{33} \\
O_{20} &= I_{20} \cdot W_{00} + I_{21} \cdot W_{10} + \boxed{I_{22} \cdot W_{20}} + I_{23} \cdot W_{30} \\
O_{21} &= I_{20} \cdot W_{01} + I_{21} \cdot W_{11} + I_{22} \cdot W_{21} + \boxed{I_{23} \cdot W_{31}} \\
O_{22} &= \boxed{I_{20} \cdot W_{02}} + I_{21} \cdot W_{12} + I_{22} \cdot W_{22} + I_{23} \cdot W_{32} \\
O_{23} &= I_{20} \cdot W_{03} + \boxed{I_{21} \cdot W_{13}} + I_{22} \cdot W_{23} + I_{23} \cdot W_{33} \\
O_{30} &= I_{30} \cdot W_{00} + \boxed{I_{31} \cdot W_{10}} + I_{32} \cdot W_{20} + I_{33} \cdot W_{30} \\
O_{31} &= I_{30} \cdot W_{01} + I_{31} \cdot W_{11} + \boxed{I_{32} \cdot W_{21}} + I_{33} \cdot W_{31} \\
O_{32} &= I_{30} \cdot W_{02} + I_{31} \cdot W_{12} + I_{32} \cdot W_{22} + \boxed{I_{33} \cdot W_{32}} \\
O_{33} &= \boxed{I_{30} \cdot W_{03}} + I_{31} \cdot W_{13} + I_{32} \cdot W_{23} + I_{33} \cdot W_{33}
\end{aligned}
$$

Example Solution for GEMM: Step2 Inter-Cluster Tile Exchanging

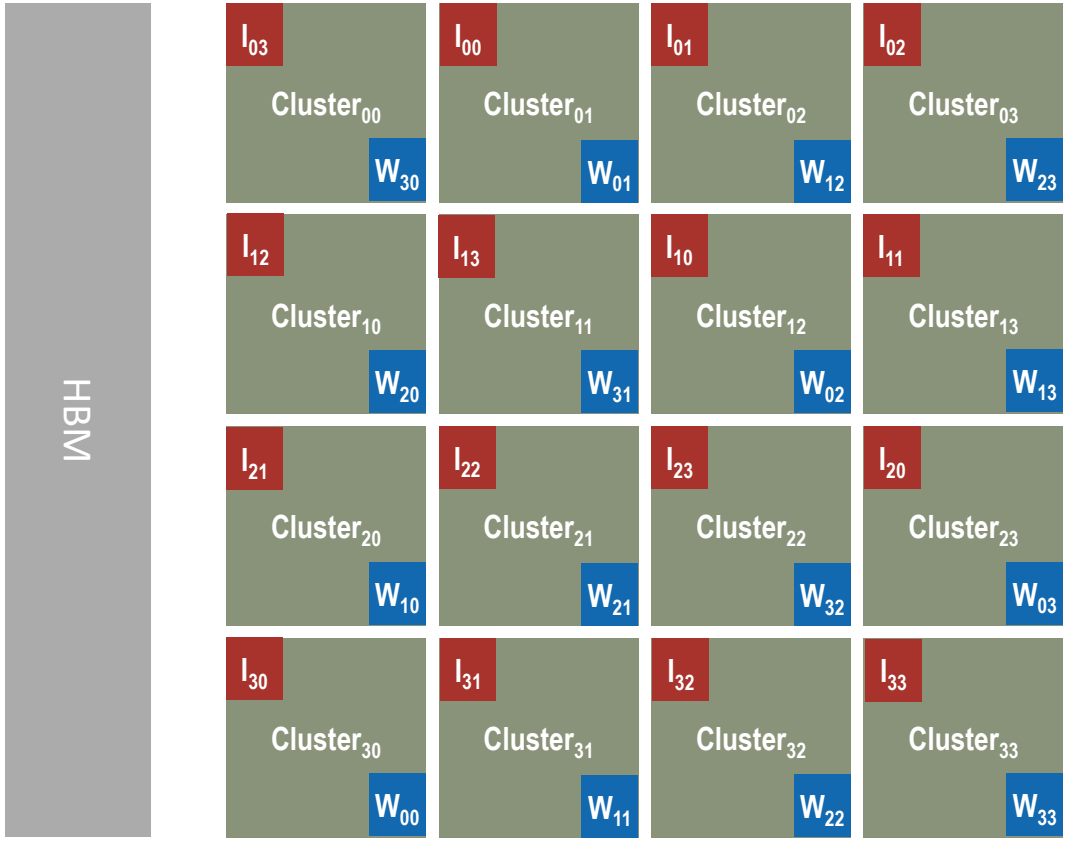

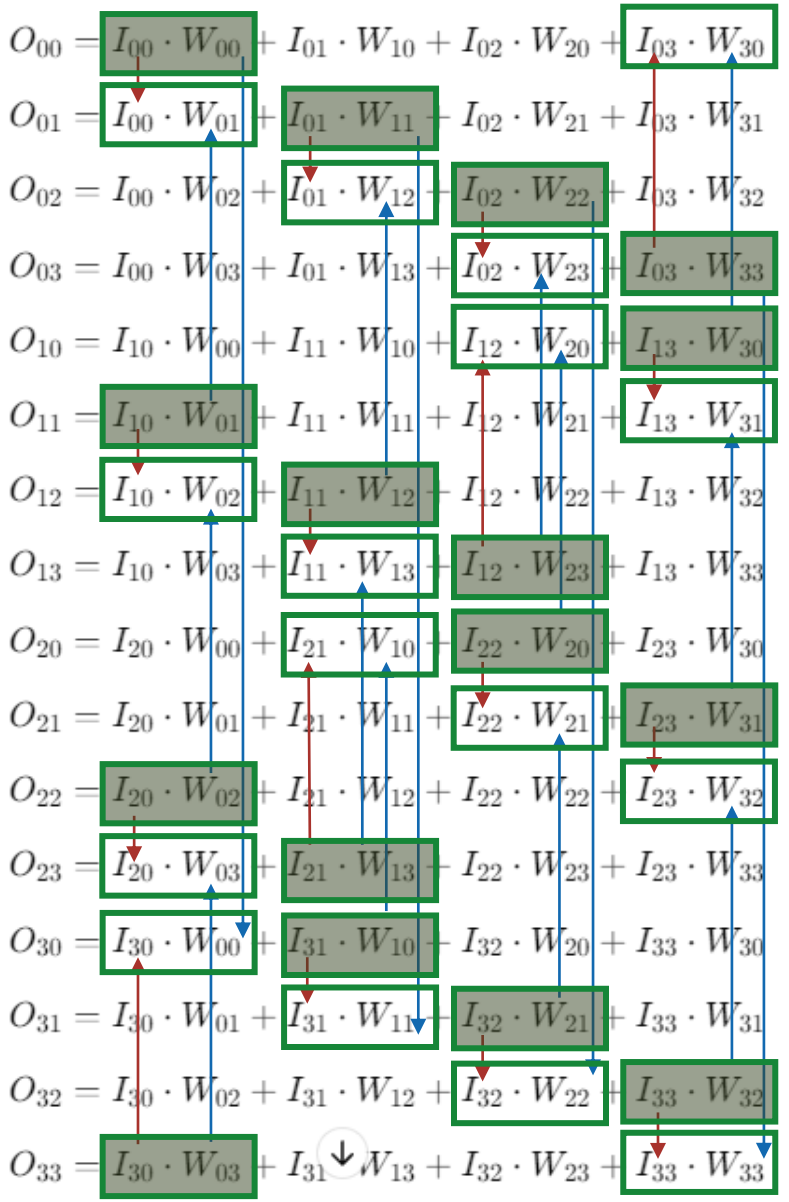

Example Solution for GEMM: Step2 Inter-Cluster Tile Exchanging

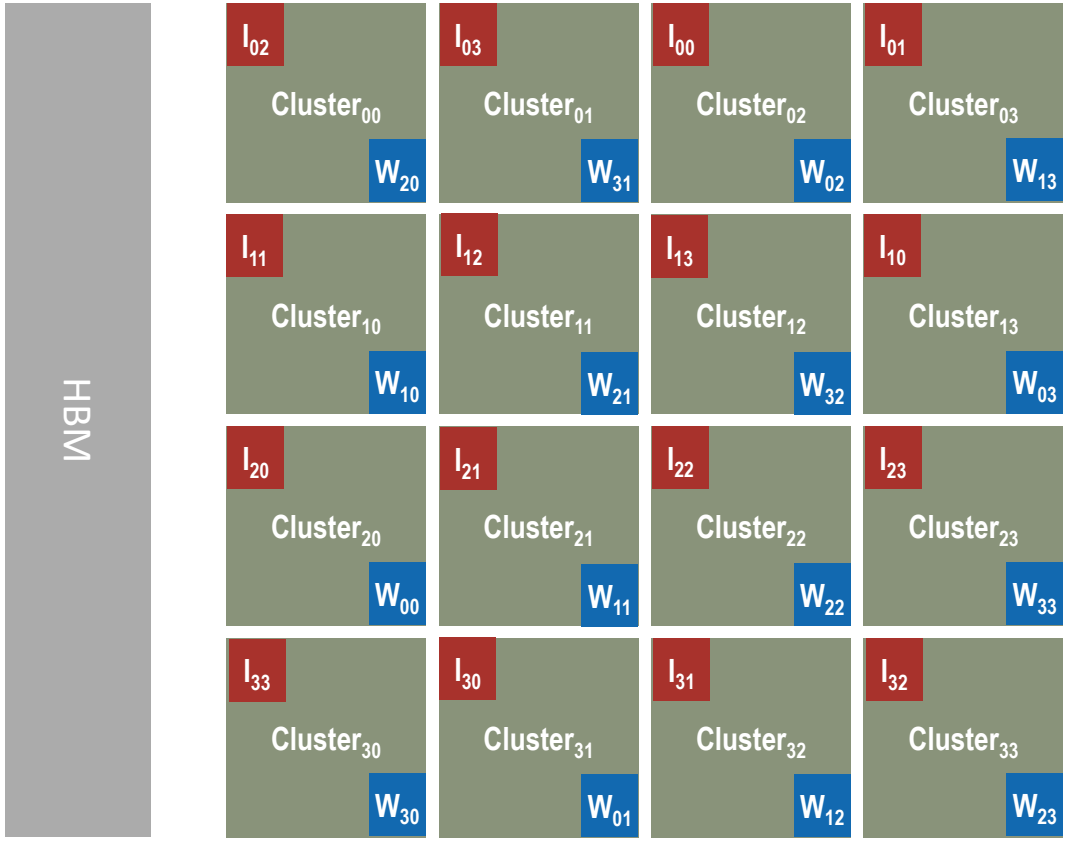

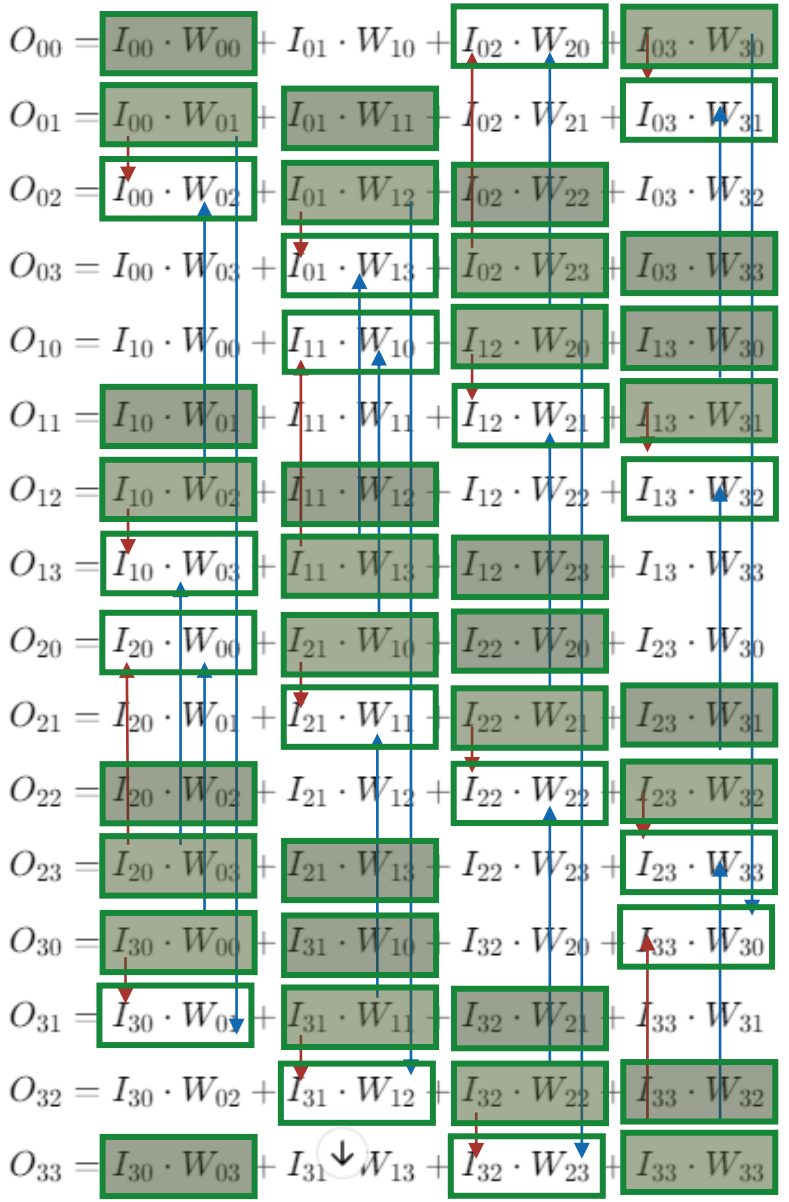

Example Solution for GEMM: Step2 Inter-Cluster Tile Exchanging

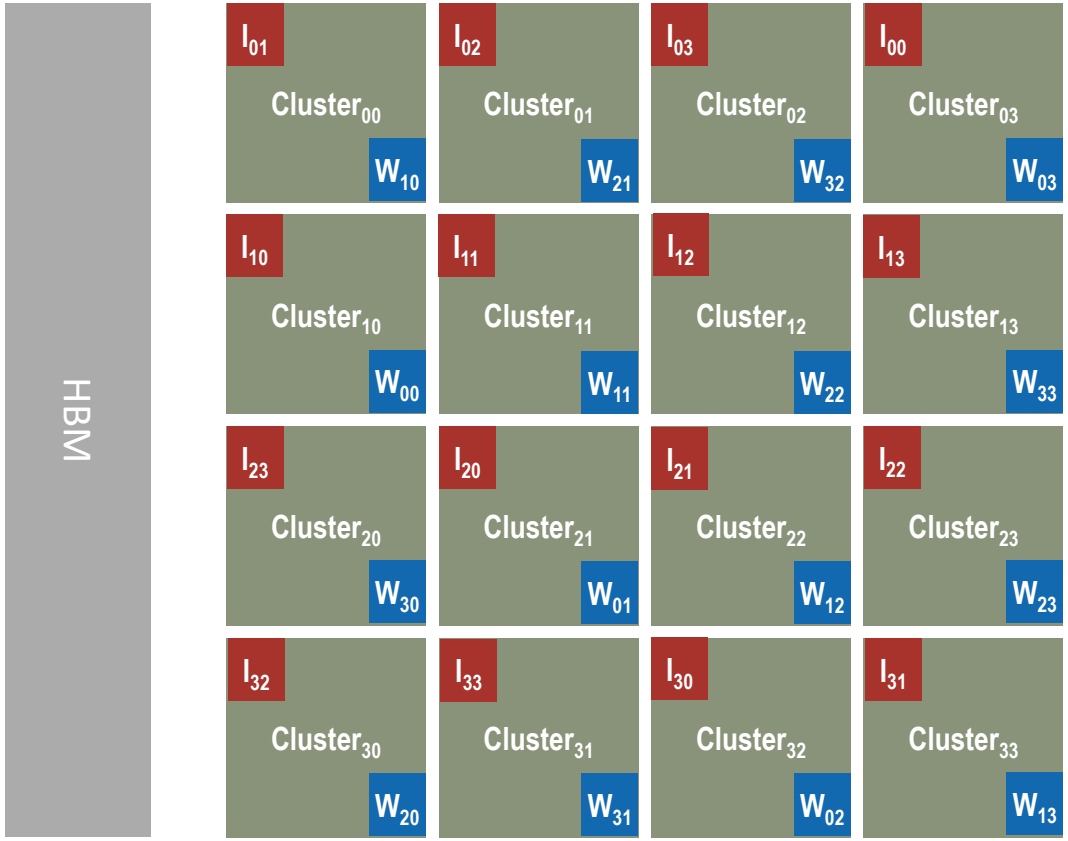

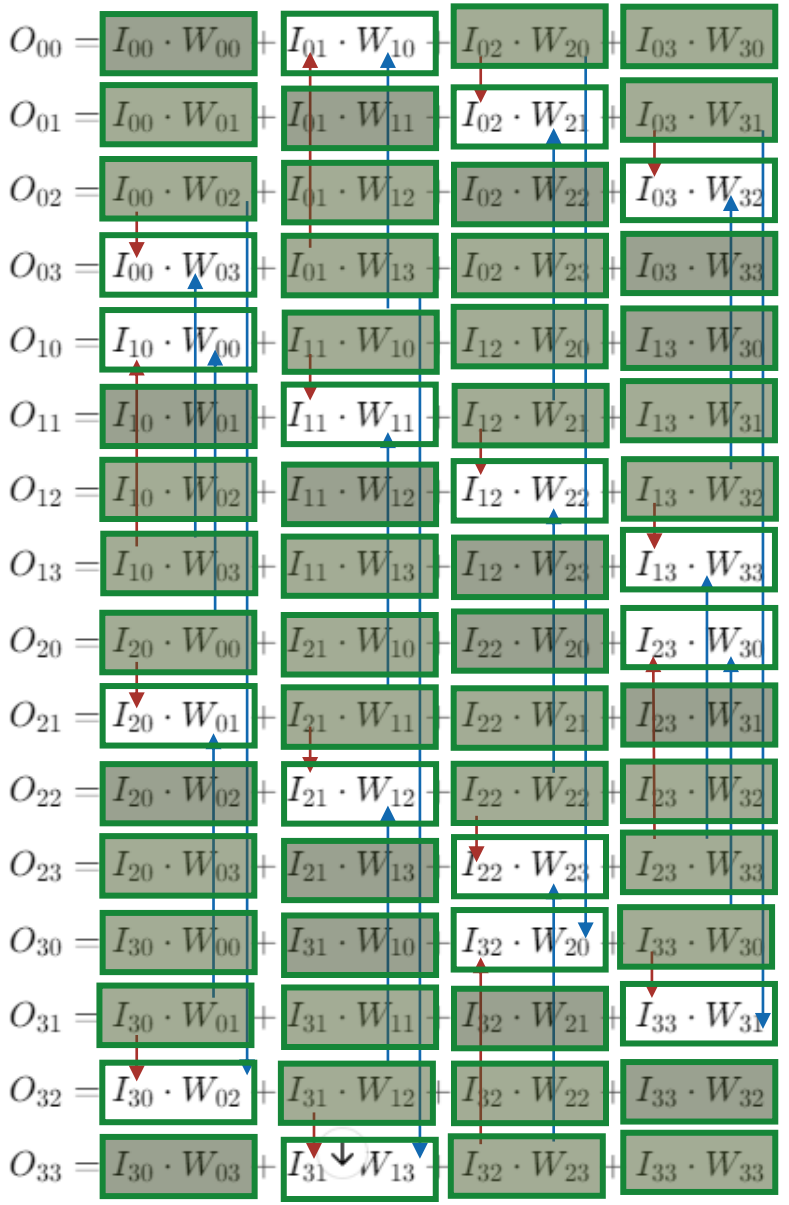

Experiment Setup: Enable Cluster-to-Cluster Comm or Not?

- RedMule
  - Each Cluster has One RedMule
  - RedMule CE array 128x32 (8TFLOPs @1GHz)
  - TCDM BW = 1024 GB/s
- HBM
  - Place HBM on left side
  - Each HBM CTRL mange 4 HBM2E channel
    - 256Byte address interleaving
    - Each HBM CTRL provide Max.BW = 205 GB/s
- NoC
  - Mesh 4x4
  - NoC link width = 2048 bits
    - Link BW = 256GB/s

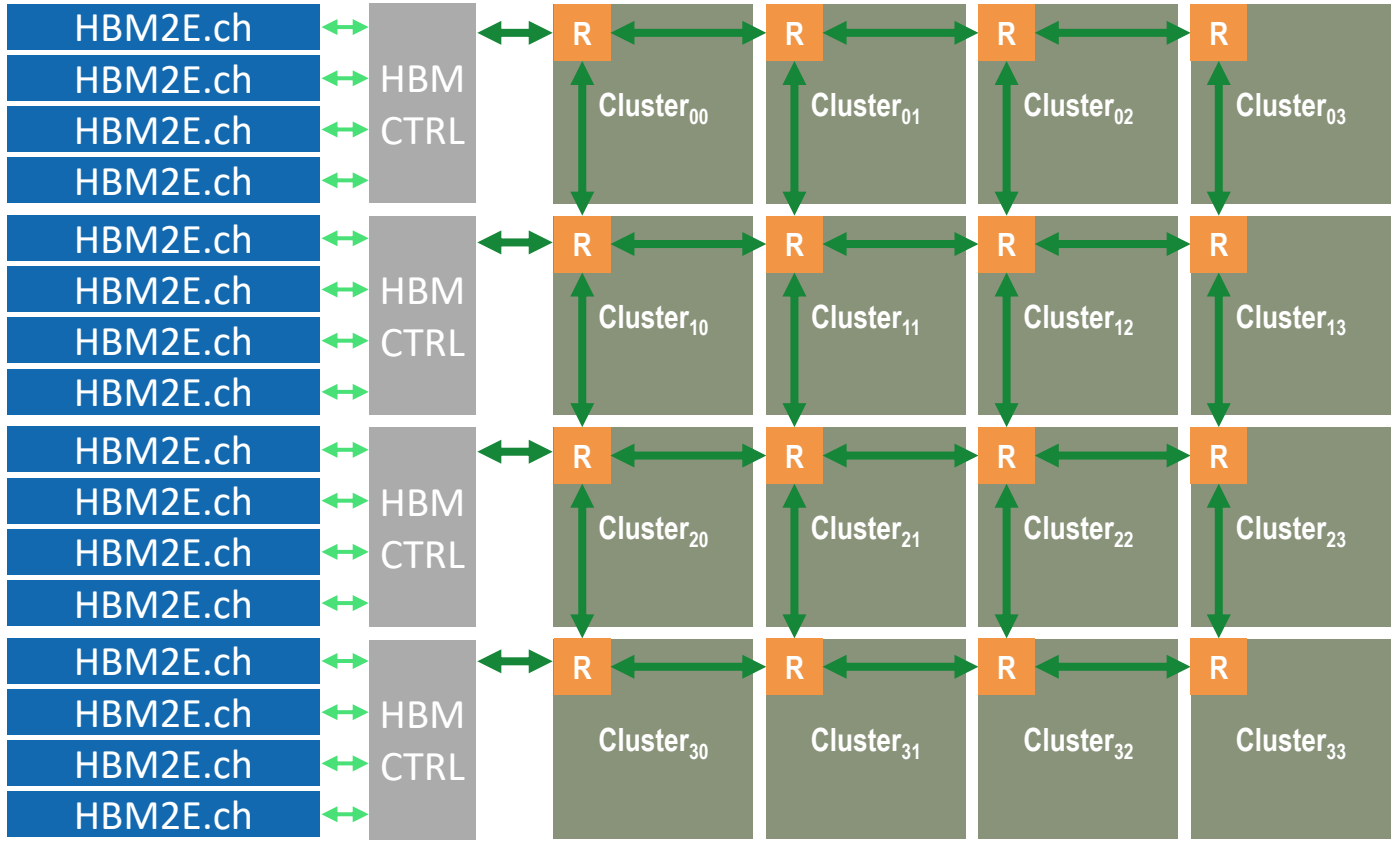

Experiment Result: Enable Cluster-to-Cluster Comm or Not?

- Benchmark GEMM
  - 16384 x16384 x16384
  - GEMM Elem Size = FP16
  - RedMule 128x32
- Varying Cluster TCDM Size
  - Tiling strategy: Max possible tile fit in TCDM
    - Tile dimension $M = N = K = \sqrt{\frac{TCDM\ Area}{5 * Elem\_Size}}$
- Results
  - No inter-cluster comm
    - HBM BW limited
    - Need 4x more TCDM size to saturate compute power
  - Enable inter-cluster comm
    - Reduce HBM traffic, better area/power efficiency

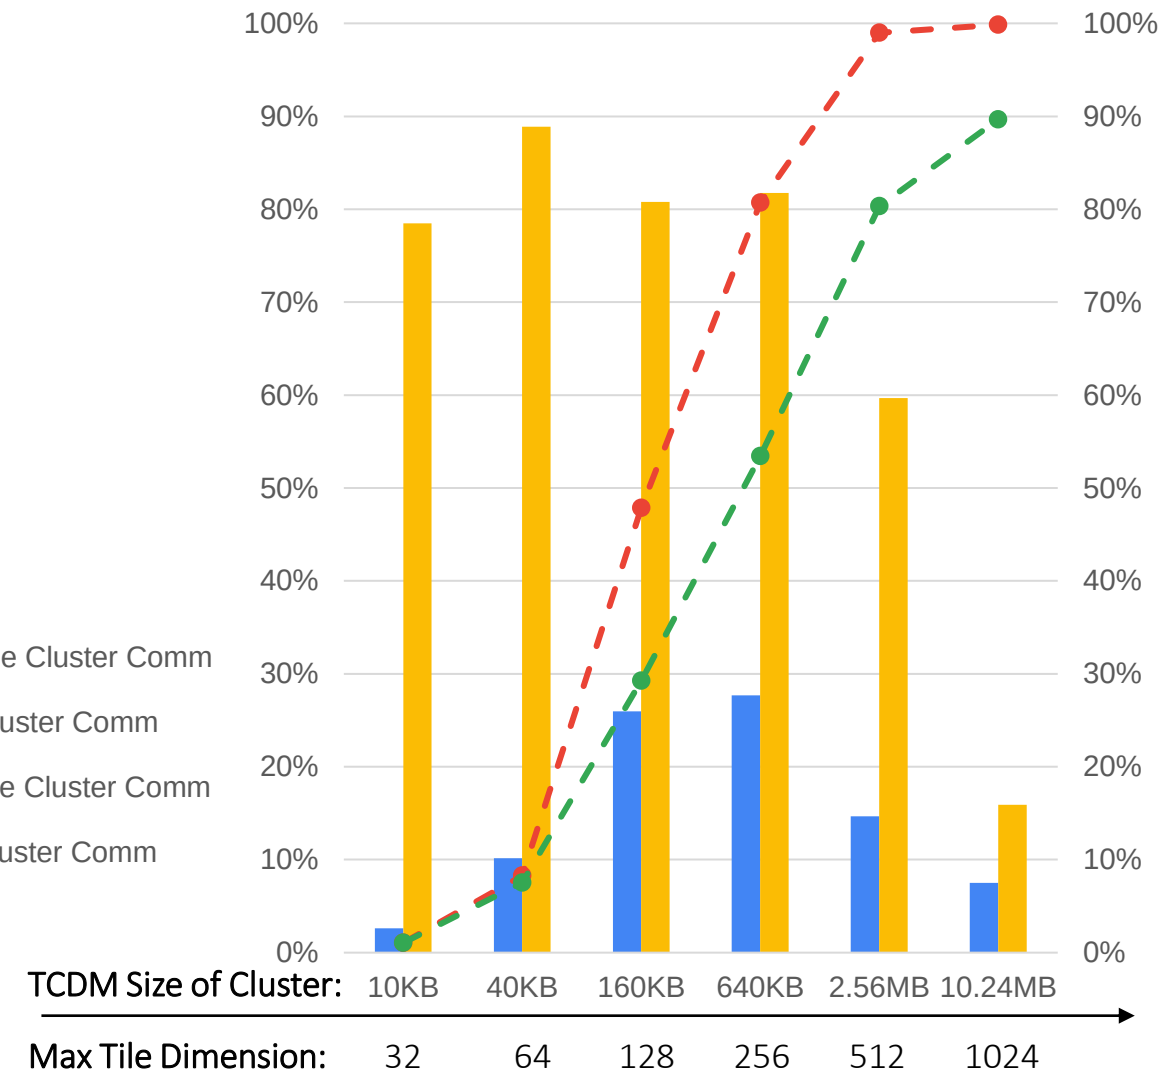

Enable Cluster-to-Cluster Comm is Scalable

- Benchmark GEMM
  - 16384 x16384 x16384
  - GEMM Elem Size = FP16
  - RedMule 128x32
- Scale-out SoftHier System
  - 4x4 Clusters + 16 HBM2E channels
  - 8x8 Clusters + 32 HBM2E channels
  - 16x16 Clusters + 64 HBM2E channels
    - 2048 TFLOPS @FP16
    - 3.2 TB/s HBM BW

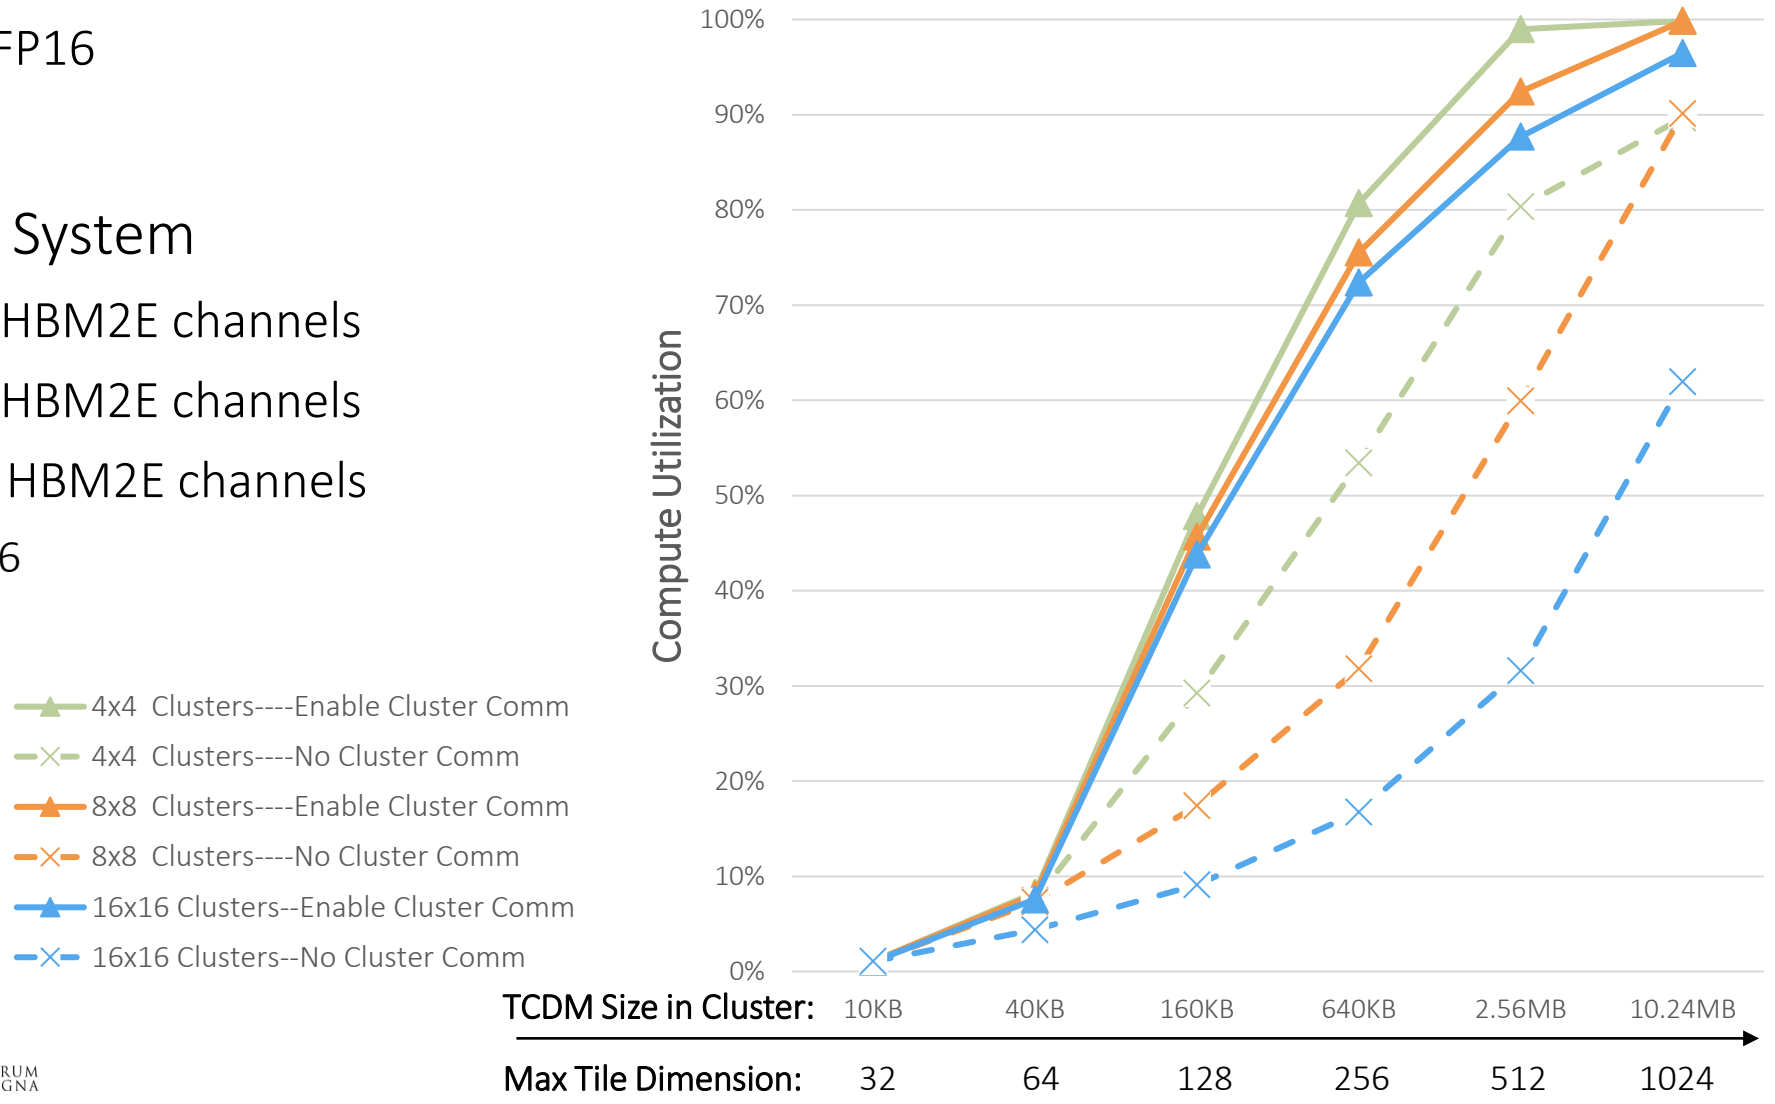

Further Discussion …

- New Questions Comes Out
  - @ RedMule optimal efficiency point
    - RedMule in cluster reaches 70% comp uti
    - But end-to-end the system shows 44% uti
    - Why? How can we optimize this?
  - Tile mapping and inter-cluster scheduling scheme
    - Is there any more schemes for large GEMM
    - How can we also leverage inter-cluster comm for MHA?
  - Inter-cluster comm vs multi-broadcasting
    - Which one is better?

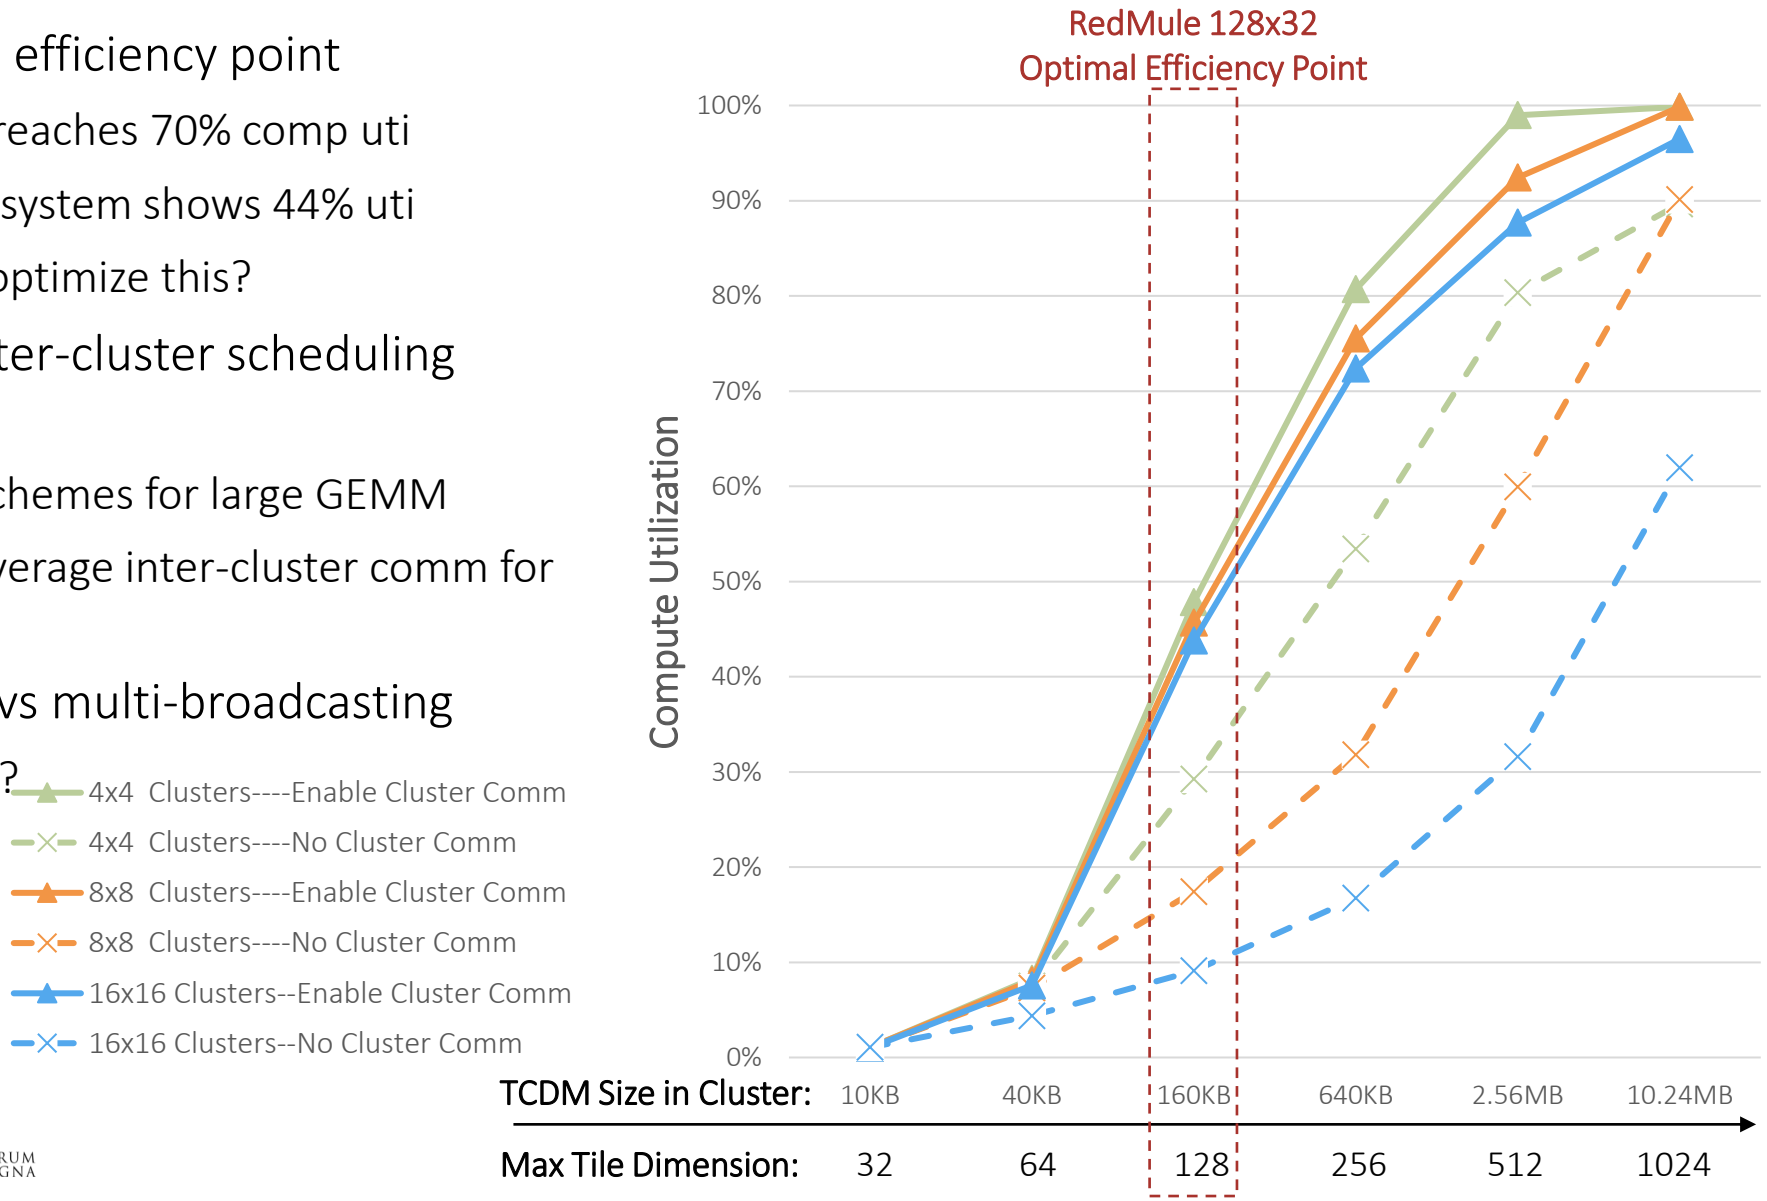

ETH zürich

ALMA MATER STUDIORUM
UNIVERSITÀ DI BOLOGNA

SoftHier Progress Update

Thomas Benz tbenz@iis.ee.ethz.ch

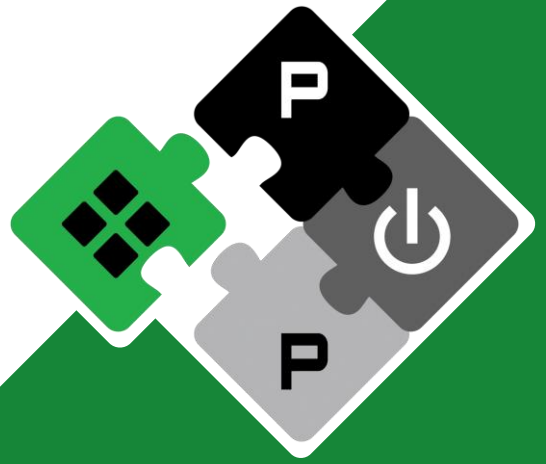

@pulp_platform

pulp-platform.org

youtube.com/pulp_platform

PULP Platform

Open Source Hardware, the way it should be!

iDMA

- Various contributions and bugfixes merged in iDMA
  - Detailed tracer will soon be enabled in Snitch (port to GVSoC pending)
  - Release early next week

- Transposition engine implemented by student overhauled
  - Support for packed SIMD types is still ongoing effort
  - Update in Snitch is ongoing

ETH zürich

ALMA MATER STUDIORUM
UNIVERSITÀ DI BOLOGNA

SoftHier Progress Update

Luca Colagrande colluca@iis.ee.ethz.ch

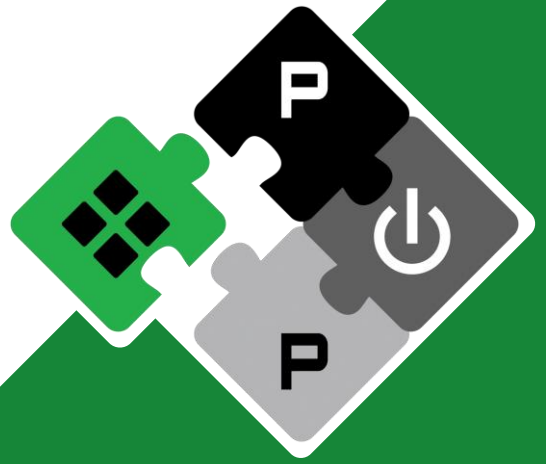

@pulp_platform

pulp-platform.org

youtube.com/pulp_platform

PULP Platform

Open Source Hardware, the way it should be!

Snitch cluster

- A lot of maintenance work:
  - Merged PRs: #165, #163, #161
  - WIP PRs: #115, #71, #158

- Implemented MHA and MLP layers for Snitch/Occamy [Link]
  - Work on streamlining data generation functions and scripts, to reuse base layer functions (e.g. GEMM, Layernorm) in composite layers (e.g. MHA, MLP)
  - Set up a proper Python package infrastructure to cross-reference these functions

- WIP on full encoder block